

Learning to find good correspondences

Kwang Moo Yi^{1,*}, Eduard Trulls^{2,*}, Yuki Ono³, Vincent Lepetit⁴, Mathieu Salzmann², Pascal Fua²

¹Visual Computing Group, University of Victoria ²Computer Vision Laboratory, École Polytechnique Fédérale de Lausanne

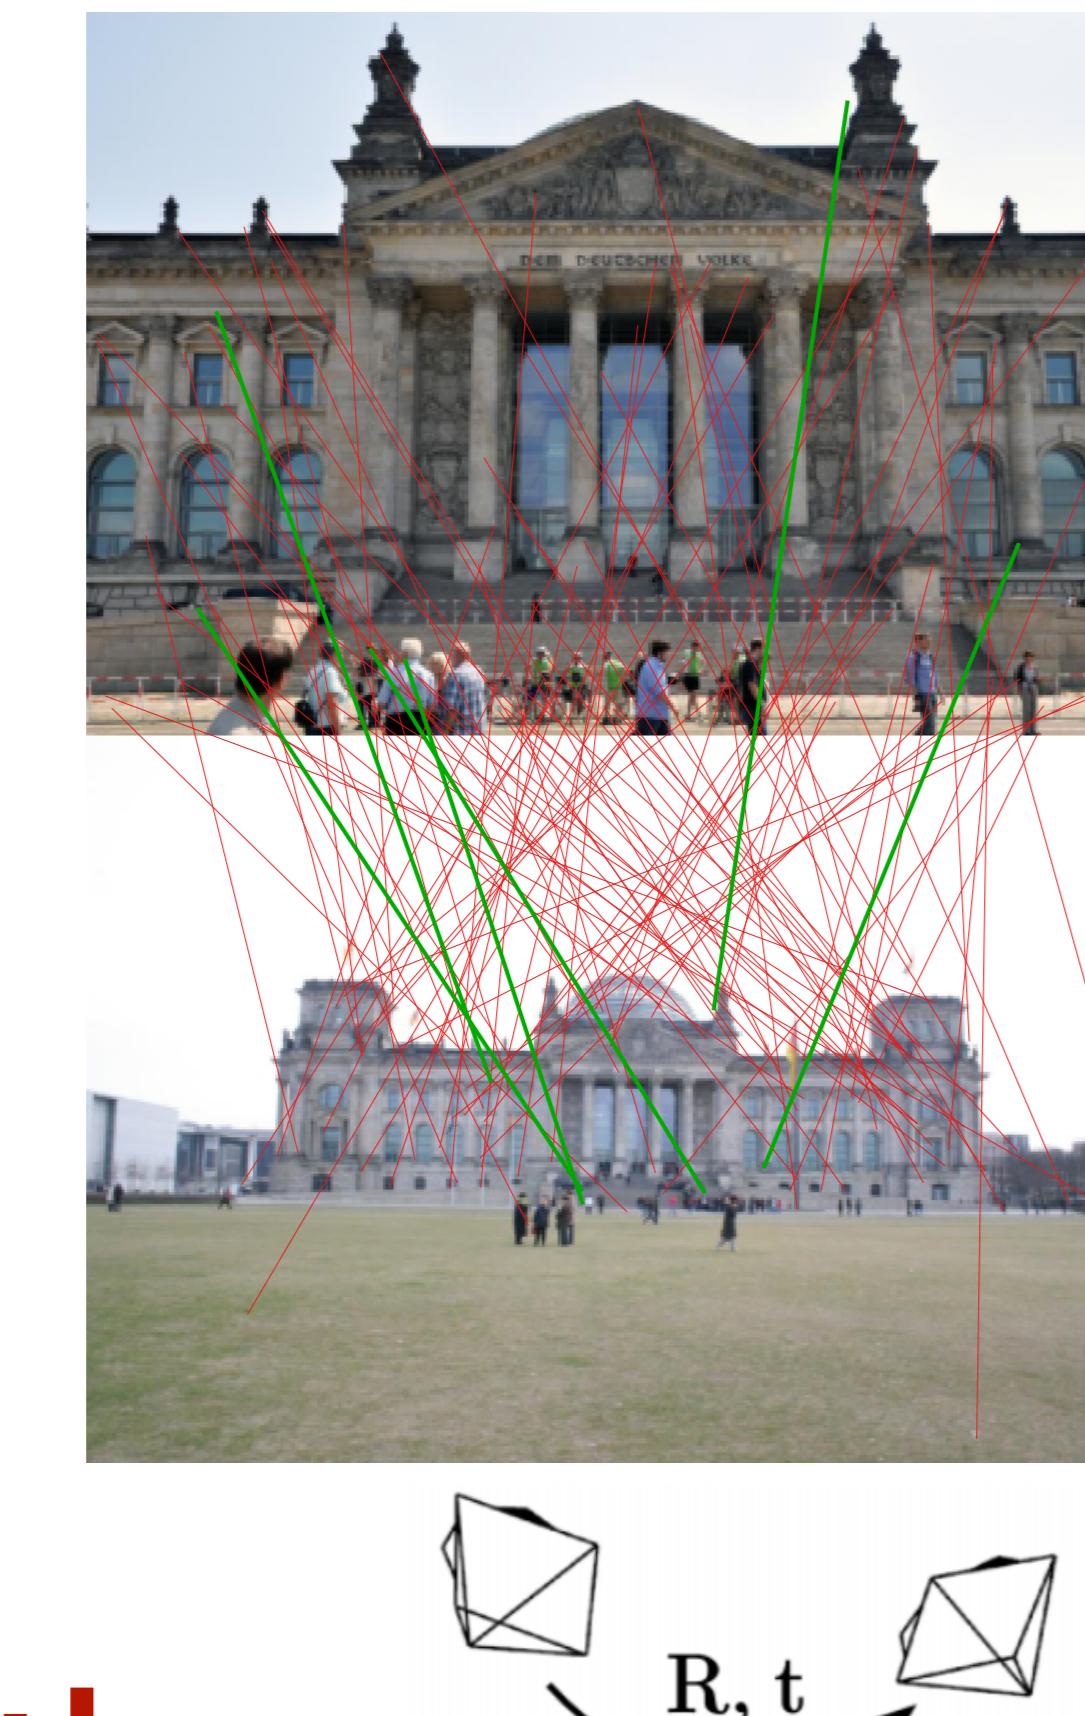
³Sony Imaging Products & Solutions Inc. ⁴Institute for Computer Graphics and Vision, Graz University of Technology (*: equal contribution)

Contributions

We solve **sparse correspondence** with **deep networks**. Classical pipeline: (a) find putative matches (e.g. SIFT); (b) find inliers (e.g. RANSAC); (c) retrieve camera motion.

Our approach:

- **Input:** correspondences. **Output:** weights.
- Unordered data → **multi-layer perceptrons**.
- **Global context** from non-parametric units.
- **Hybrid loss:** joint classification & regression.
- **State of the art** results on indoors/outdoors.



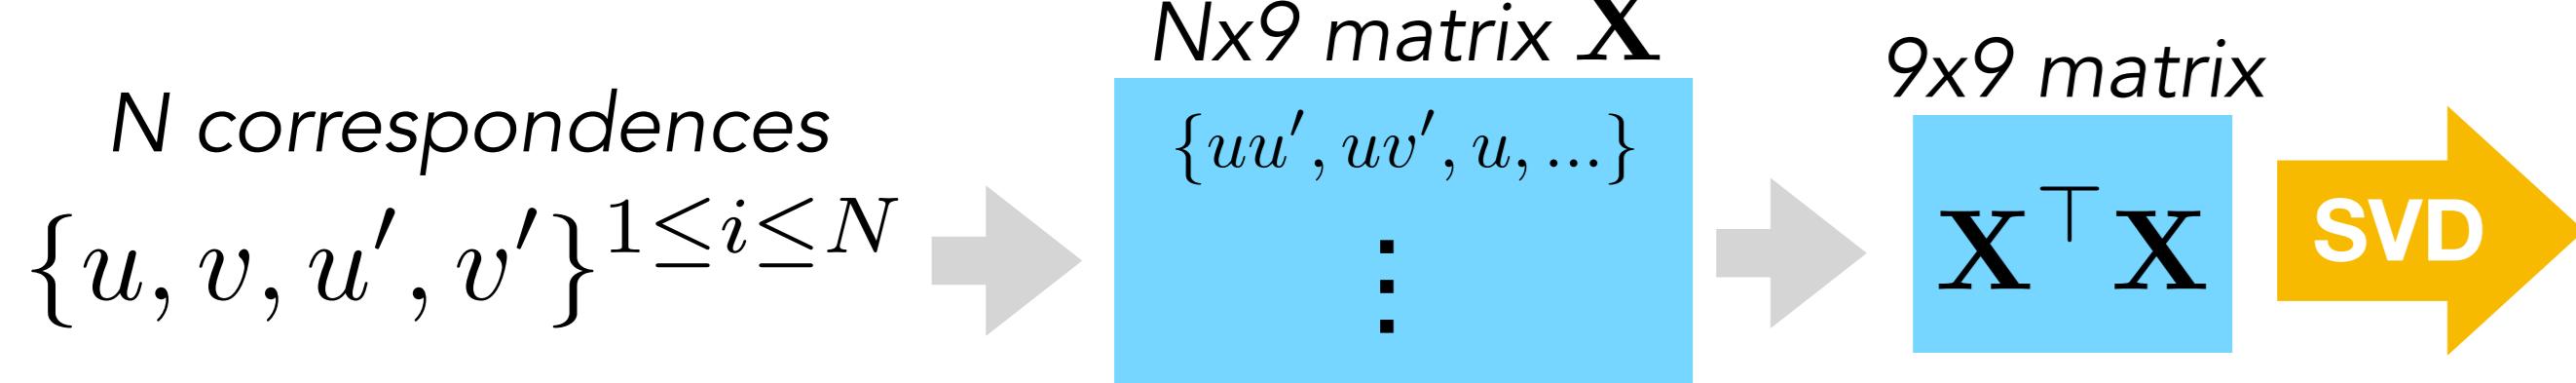
Collecting the Ground truth



Can't have pixel-to-pixel correspondences. We propose to use **only the pose as ground truth**. We can recover it with off-the-shelf SfM.

Learning with regression

- 8-point algorithm: closed-form solution for the Essential matrix:



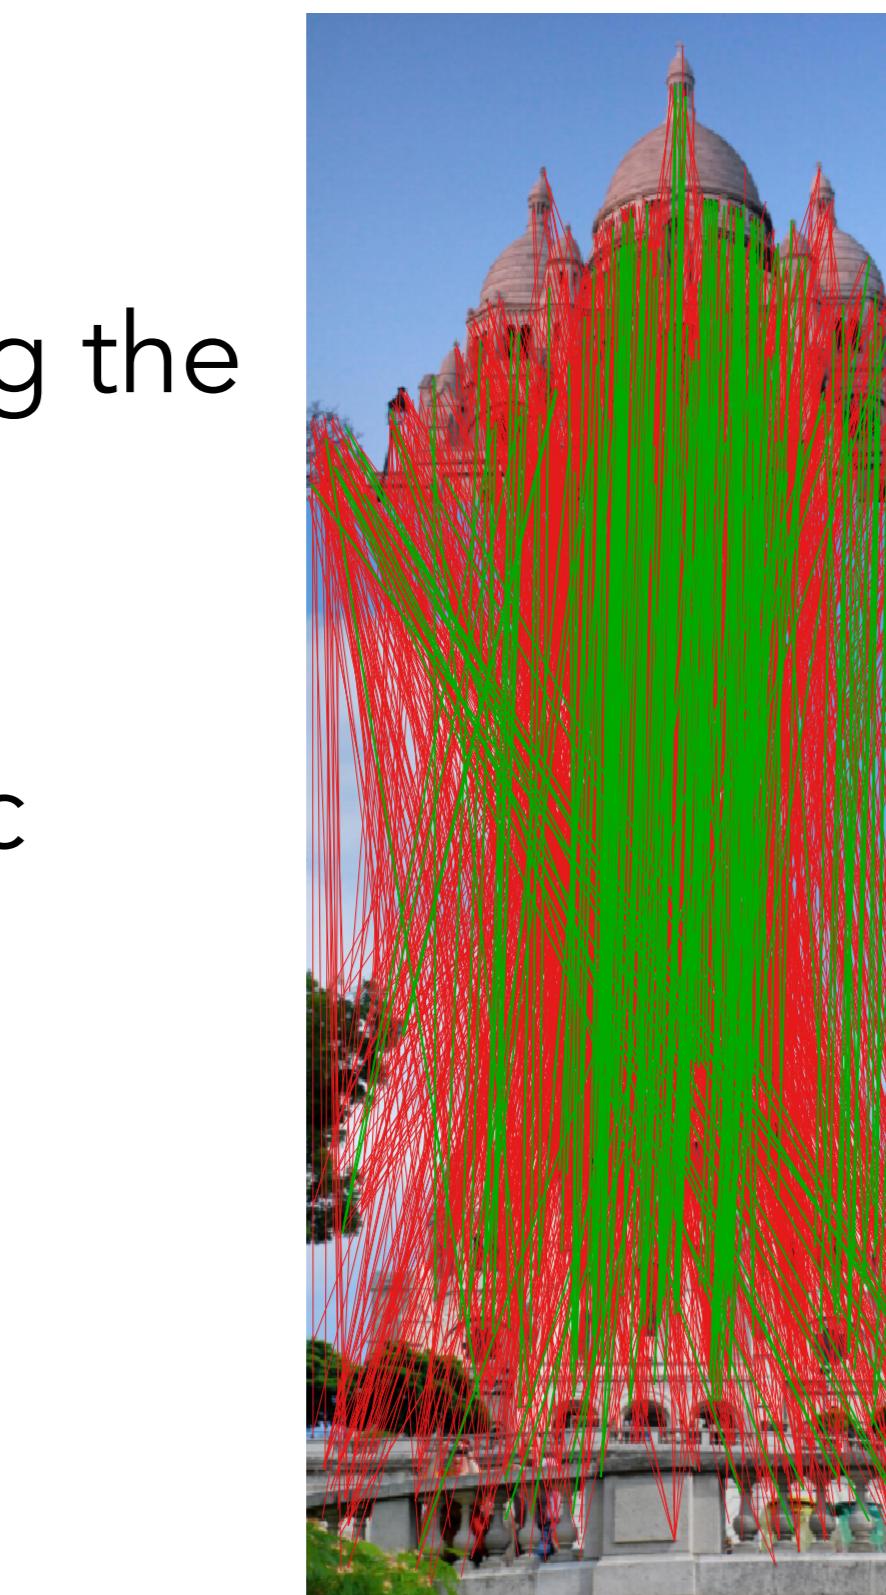
- Problem: **weak to outliers**. Solution: **weighted 8-point**, using the weights from the network: $\mathbf{X}^\top \text{diag}(\mathbf{w}) \mathbf{X}$. Fully differentiable.

$$\mathcal{L}_e(\Phi, \mathbf{x}_k) = \min \left\{ \| \mathbf{E}_k^* \pm g(\mathbf{x}_k, \mathbf{w}_k) \|^2 \right\}$$

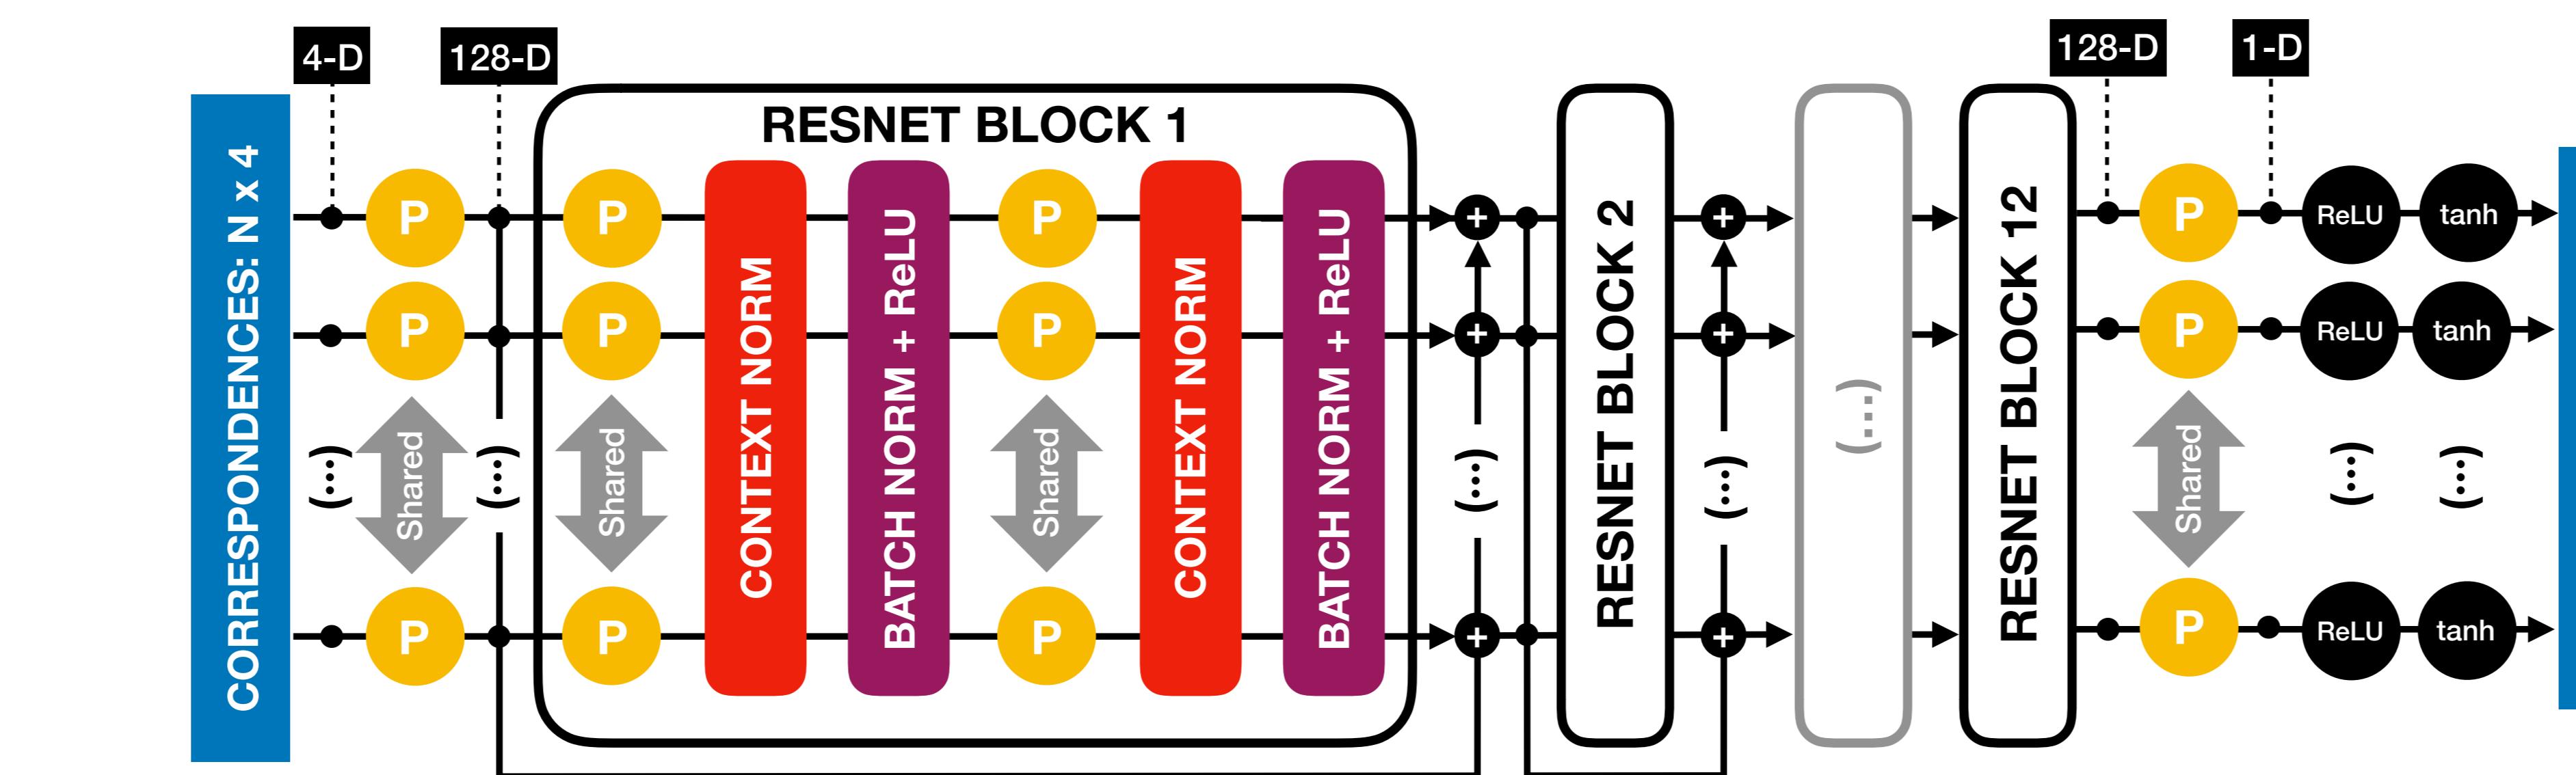
Learning with classification

- Learning outlier rejection implicitly by regressing the pose is too hard. **Network does not converge**.
- Solution: create **training labels from epipolar constraints**. How? threshold over the symmetric epipolar distance. Noisy but good enough!
- Loss: standard binary cross-entropy.

$$\mathcal{L}_x(\Phi, \mathbf{x}_k) = \frac{1}{N} \sum_{i=1}^N \gamma_k^i H(y_k^i, S(o_k^i))$$



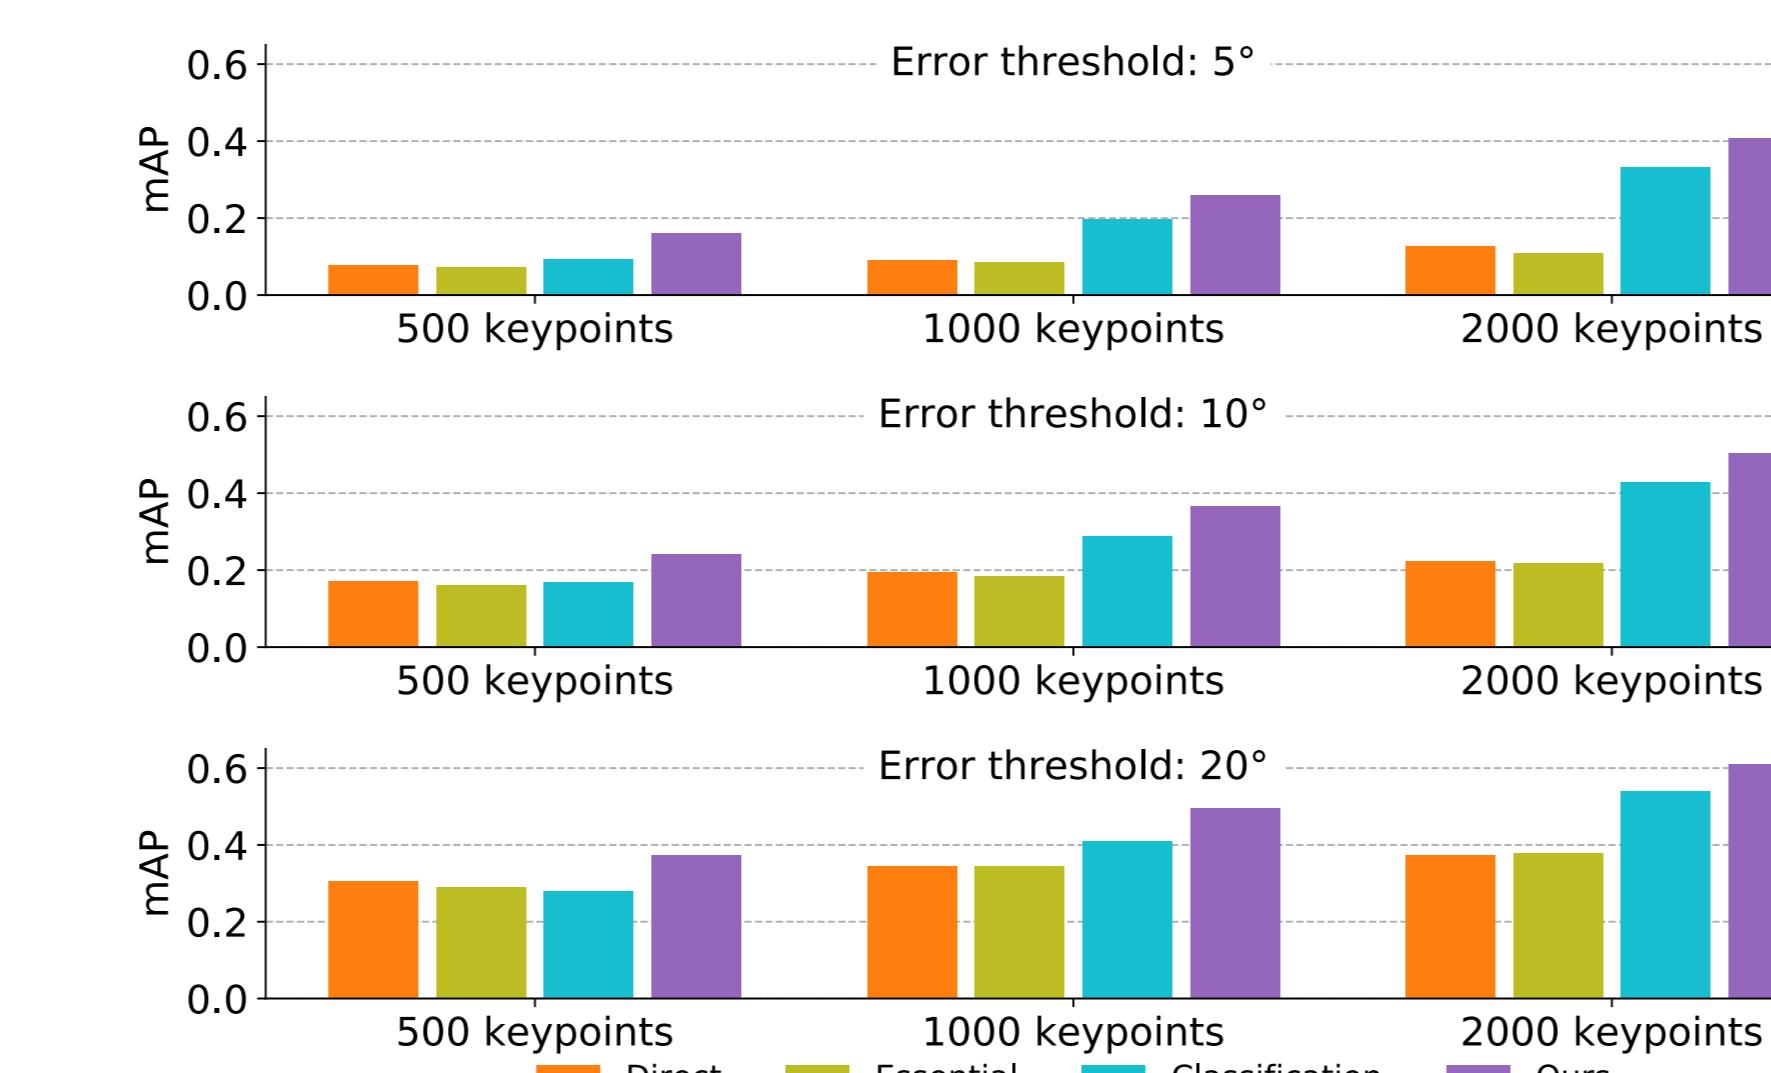
Outlier Rejection Network



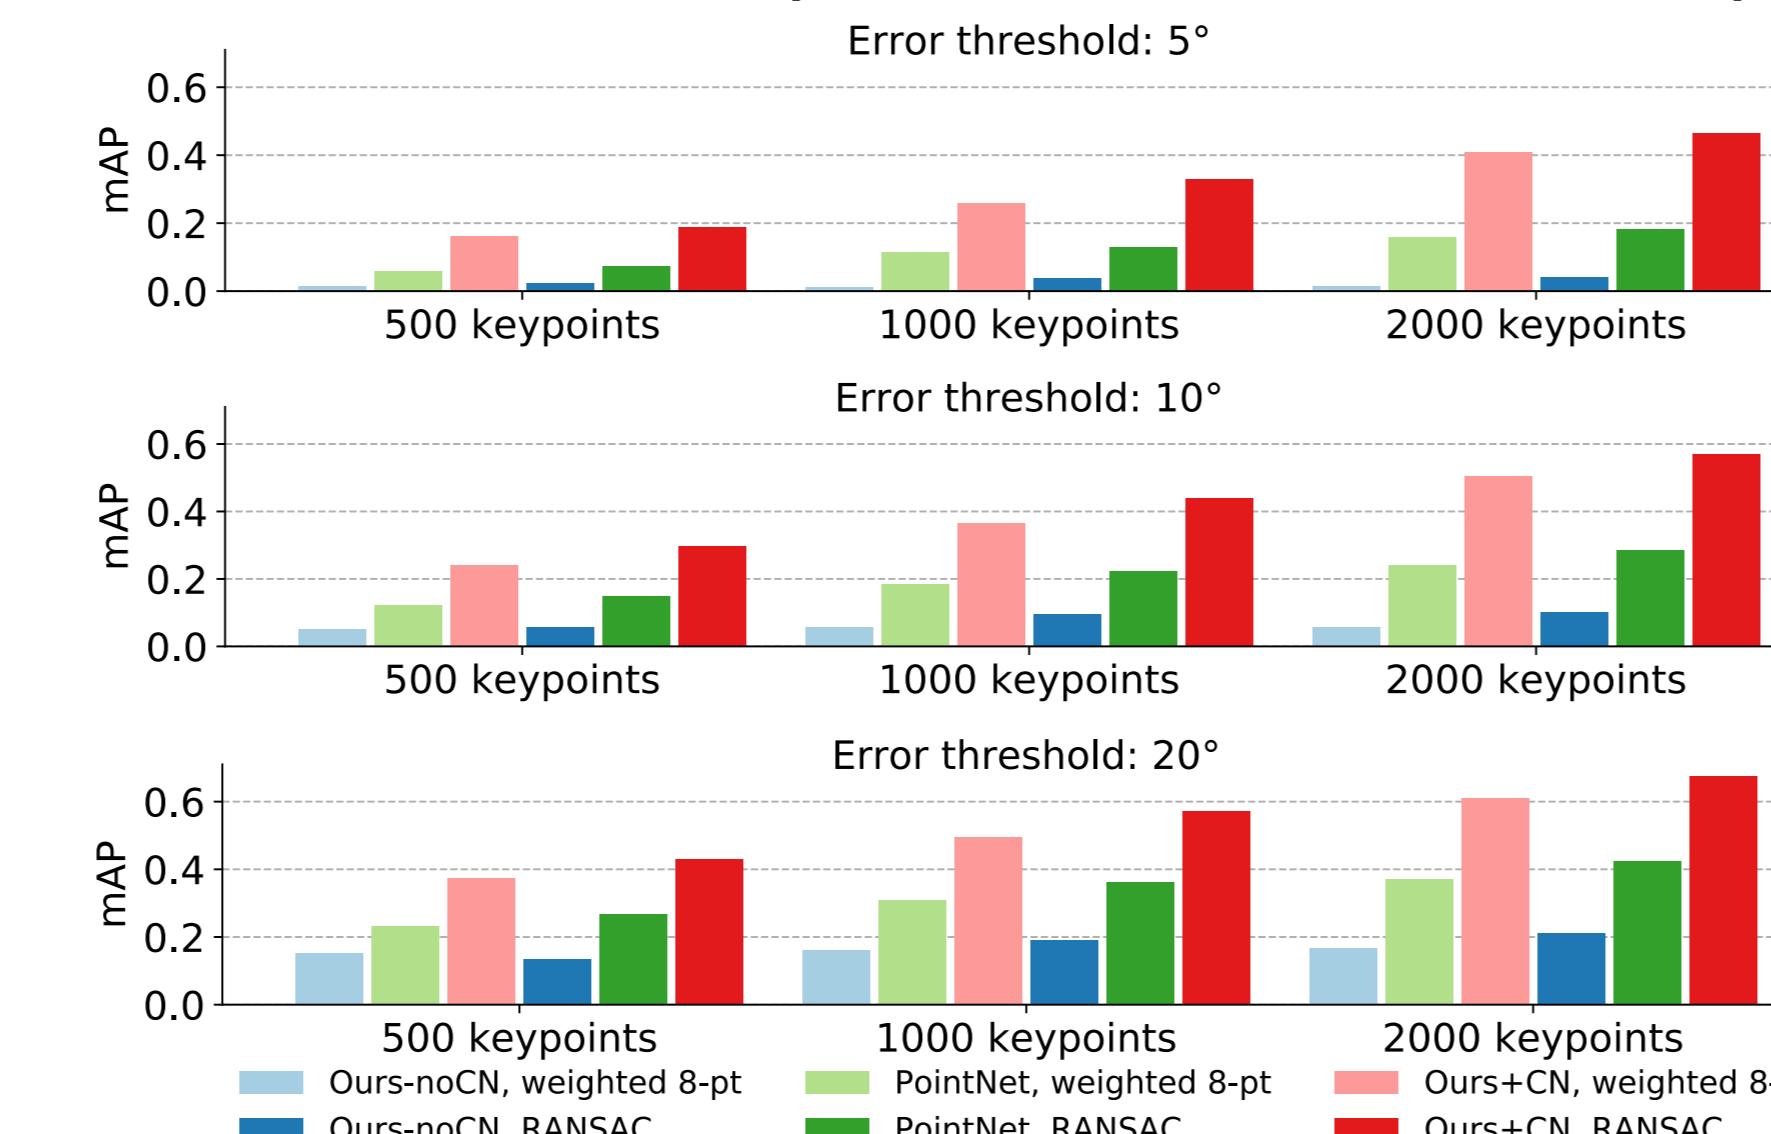
- **Input:** N correspondences $\{u, v, u', v'\}^{1 \leq i \leq N}$. **Output:** N weights.
- **Problem:** input data is unordered. Output should be permutation-invariant. Not feasible with e.g. convolutional or fully-connected layers.
- **Solution (PointNet):** Multi-layer, weight-sharing perceptrons.
- Deep network: 12 resnet-style blocks. Still **very small!**
- Each point is processed individually! We need contextual information. PointNet solution: **global feature**, pooled with a second network.
- **Our solution:** embed into the feature maps with **Context Normalization**.

Ablation: Loss & Context

- Classification required to converge. Hybrid loss does best: $\mathcal{L}(\Phi) = \sum_{k=1}^P (\alpha \mathcal{L}_x(\Phi, \mathbf{x}_k) + \beta \mathcal{L}_e(\Phi, \mathbf{x}_k))$



- PointNet-style context works, but our simple Context Norm outperforms it on this problem.



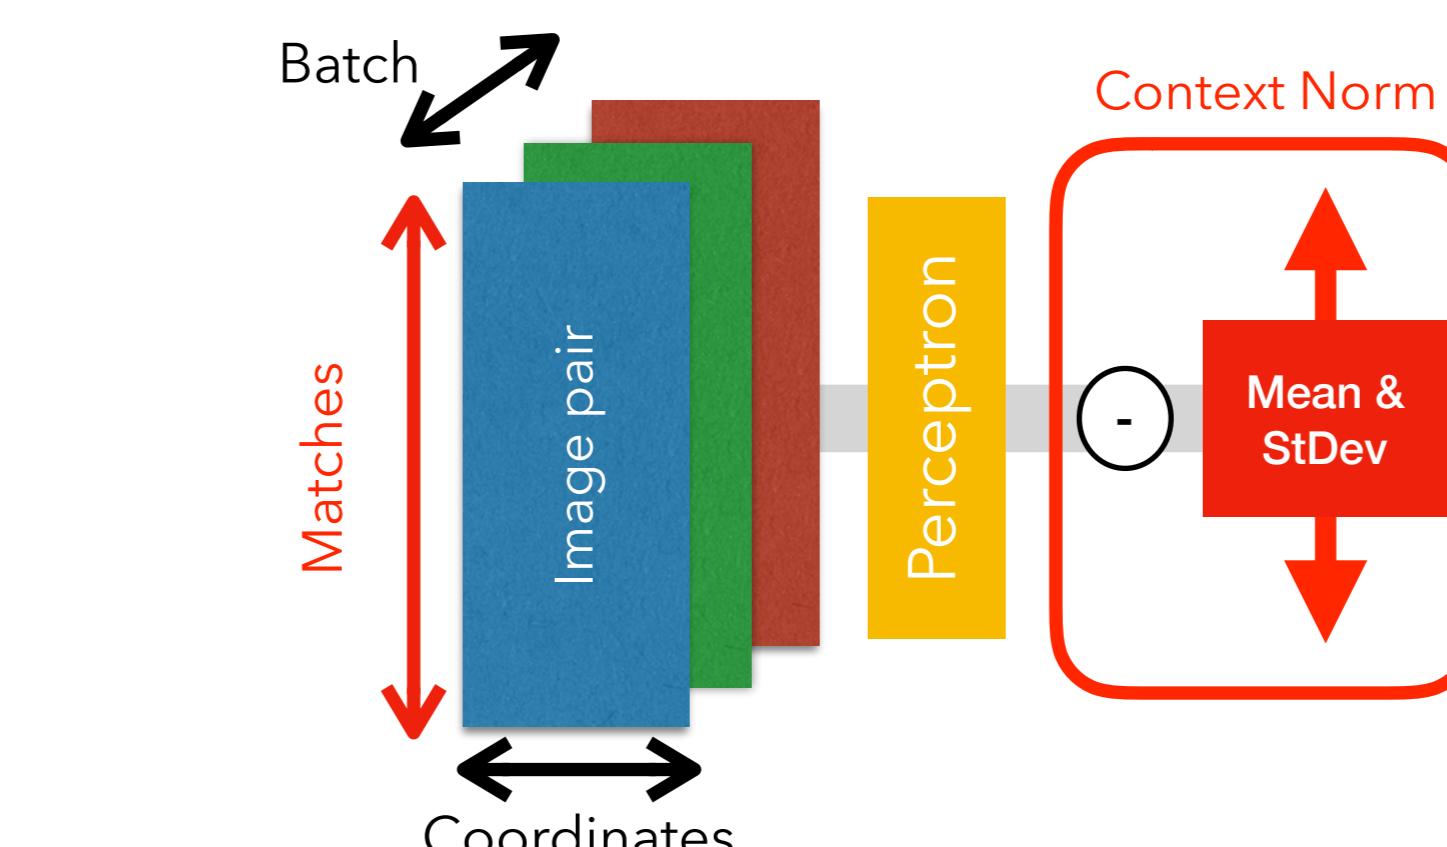
Context Normalization

- Simple, non-parametric normalization. Given features \mathbf{o}_i^l at layer l :

$$CN(\mathbf{o}_i^l) = \frac{(\mathbf{o}_i^l - \mu^l)}{\sigma^l}$$

$$\mu^l = \frac{1}{N} \sum_{i=1}^N \mathbf{o}_i^l, \quad \sigma^l = \sqrt{\frac{1}{N} \sum_{i=1}^N (\mathbf{o}_i^l - \mu^l)^2}$$

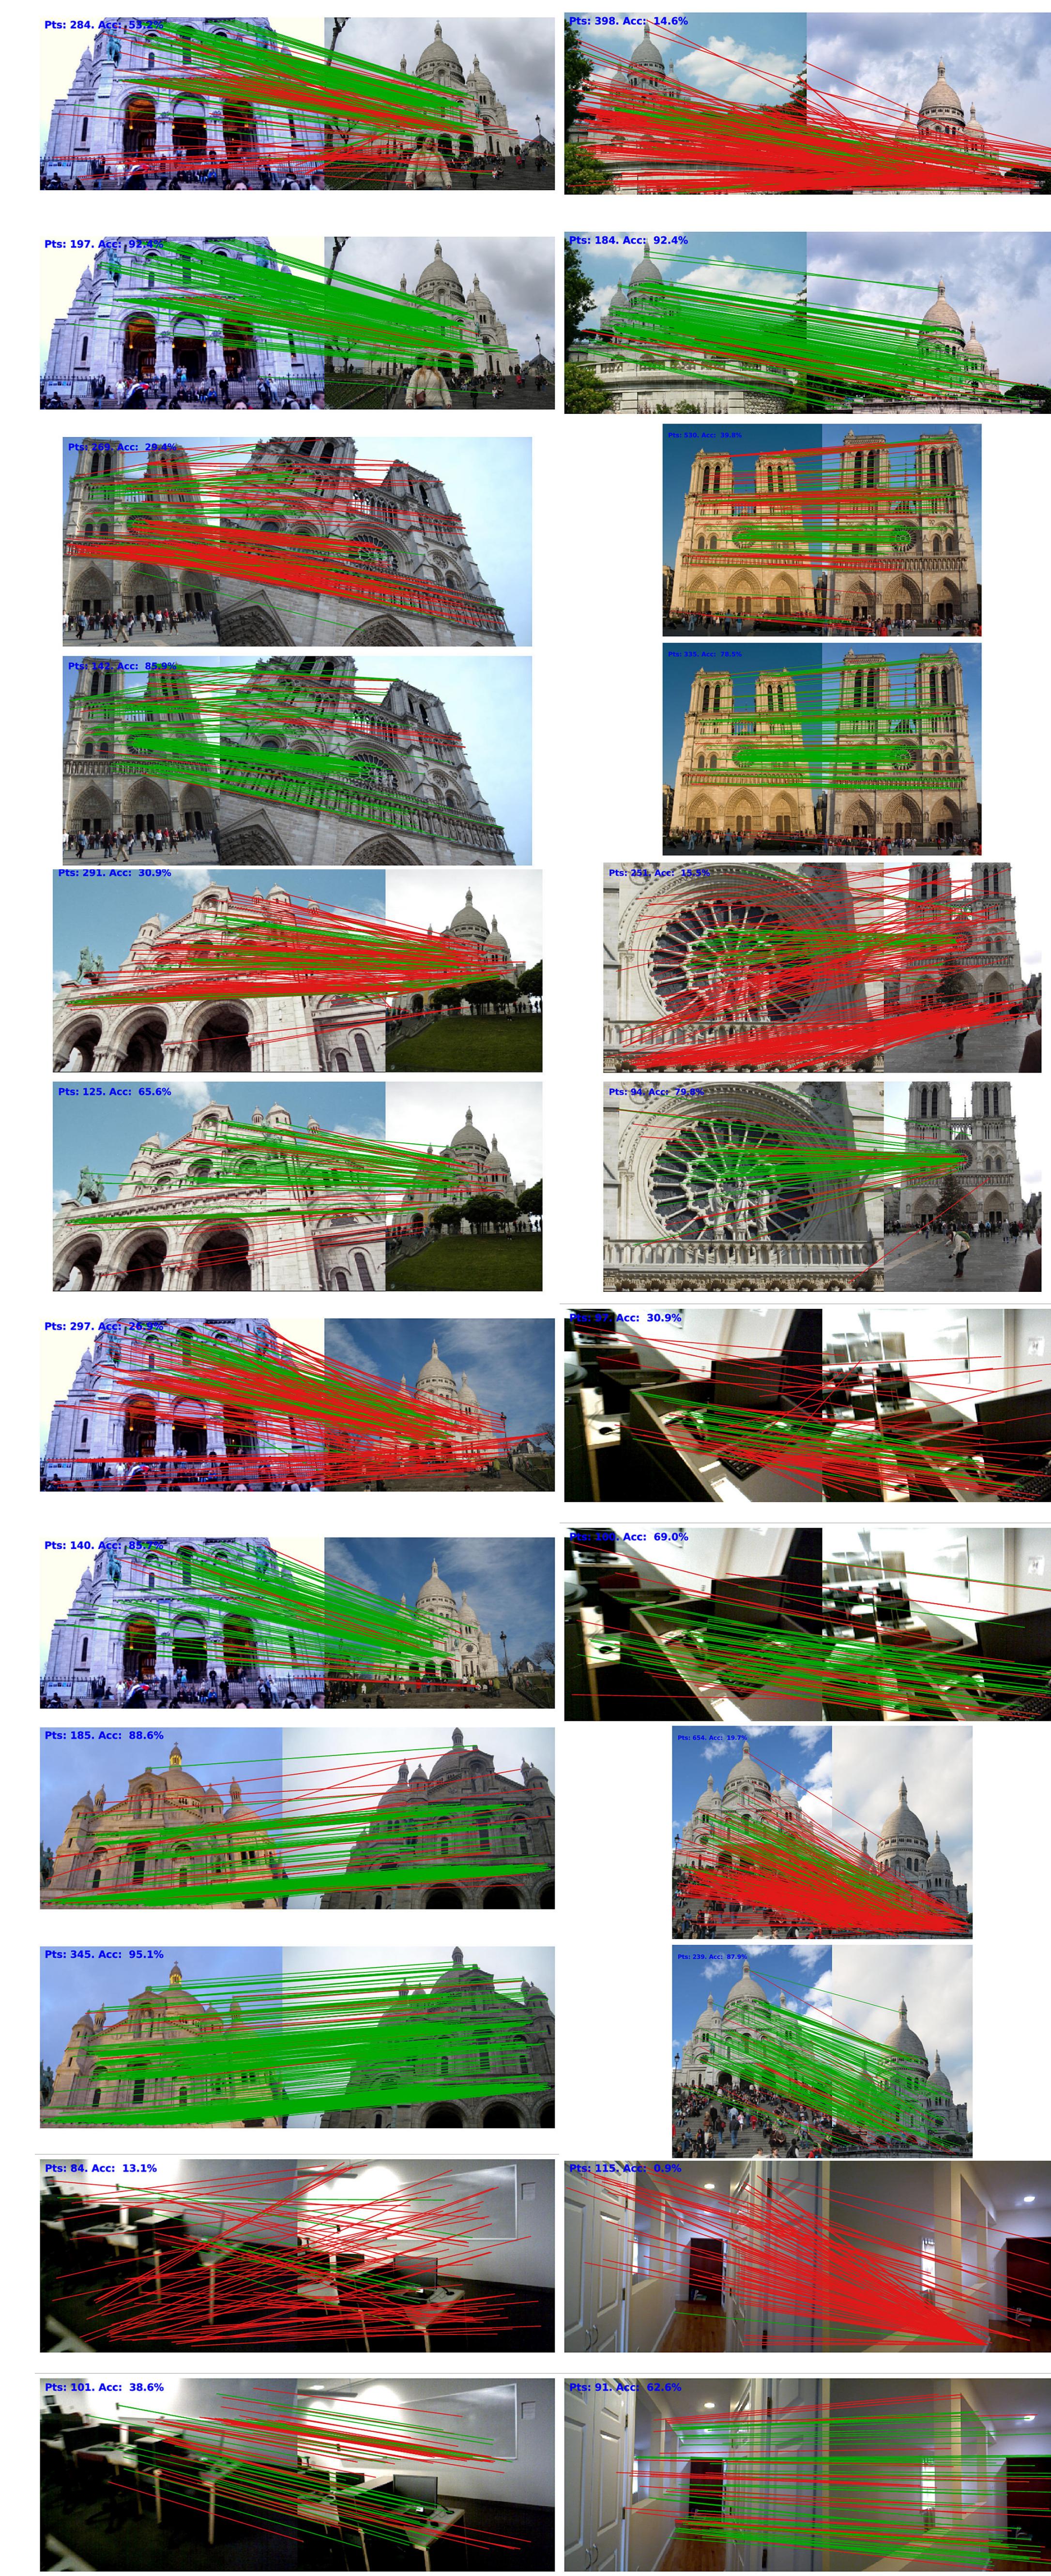
- Similar to BN/LN, but nothing is learned. Same operation for training/inference.
- Operates separately over image pairs:



- In image stylization: **Instance Norm**.

Qualitative results

- Top: **RANSAC**. Bottom: **Ours**. Same input.
- Drawing **inliers only**. Pictured in **green** if they are below the ground truth epipolar distance threshold, and in **red** otherwise.



Evaluation

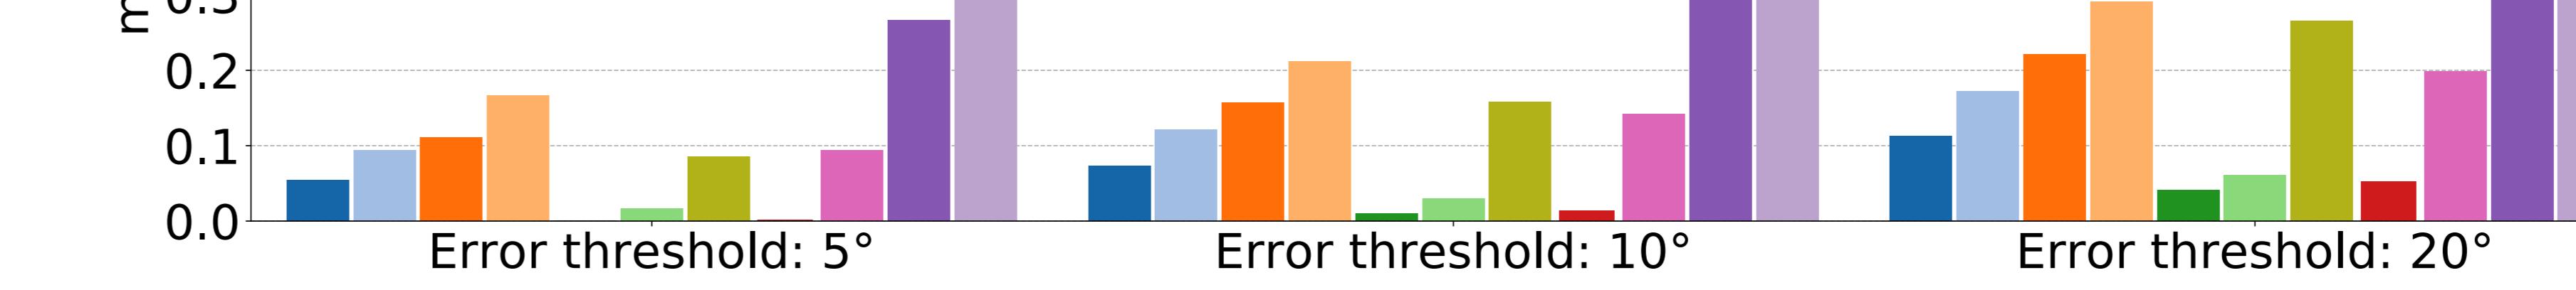
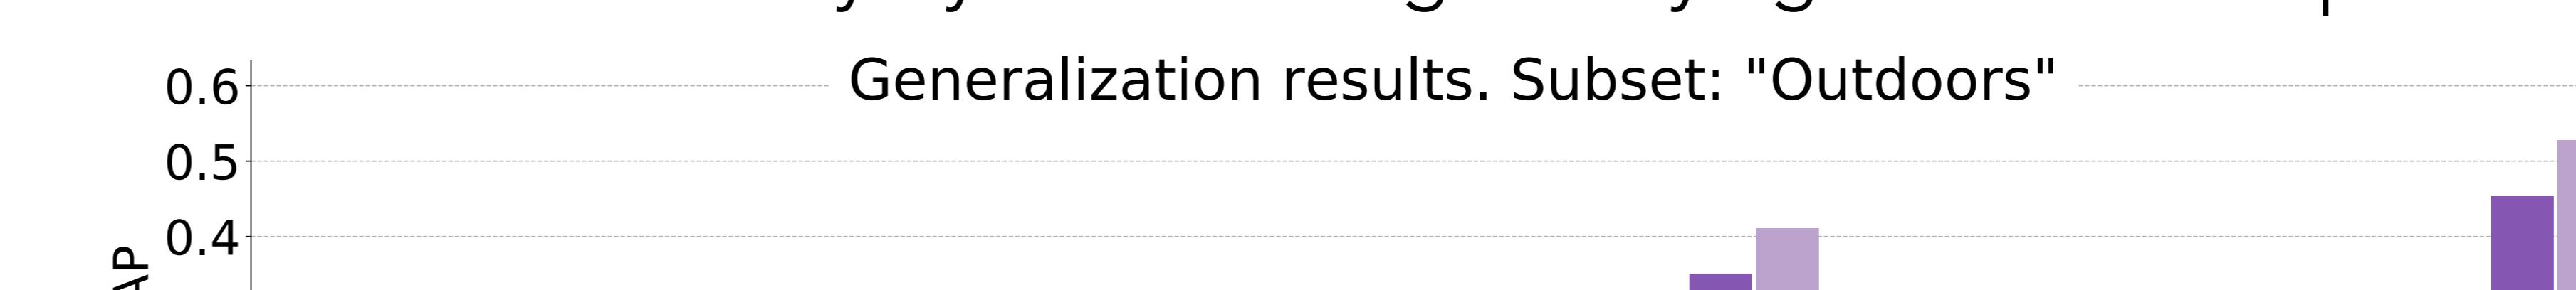
- **Datasets:** indoors (SUN3D) and outdoors (YFCC100+SfM).

- Our models are trained on a **single sequence from each**.

- **Baselines:** sparse (RANSAC variants, GMS) & dense (G3DR, DeMoN).

- **Metric:** angular error between ground-truth & estimated R/T.

Determine accuracy by thresholding at varying values & compute mAP.



- **Outdoors:** great improvements. **Indoors:** still better than dense SoA.
- For testing we do not need differentiability! We run **ours (one forward pass) then RANSAC**. Improves performance (2x) and speed (17x!).